

## Assignment - 4

Name :- Seihti Mishra

Roll no :- 2301410005

Course: Btech CSE (Cyber Security)

Ques1:- Explain race conditions with a real-world example outside of computing, and show how mutual exclusion addresses it.

Ans:- Imagine two people trying to withdraw money from the same bank account at the same ATM at the same time. Both see the balance as 500. Each withdraws ₹500, and the system allows both, causing inconsistency. The final outcome depends on "who finishes first" not on correct logic → This is a race condition.

→ How mutual exclusion solves it :

A mutex would allow only one user to access/update the account at a time. When one withdrawal is in progress, the other must wait. This prevents conflicting updates and ensures consistent balance.

Ques2:- Compare peterson's solution and semaphores in terms of implementation complexity and hardware dependency.

Ans :- Peterson's solutions:-

- Pure software-level algorithm; Simple but error-prone; limited to 2 processes.
- Requires strict assumptions about atomic read/write ordering → not reliable on modern multi-core CPUs.

→ Semaphores :-

- More general:- easier for programmers; widely supported.
- Depends on hardware atomic operations like test-and-set; explicitly supported by processors.

Ques3. The producer-consumer problem can be solved using either semaphores or monitors. Identify one advantage of using monitors in a multi-core system.

Ans:- Monitors automatically combine mutual exclusion with Condition Variables, so synchronization is handled implicitly by the language/runtime.

Advantage:- They prevent low-level race conditions automatically, allowing threads on multiple cores to safely access shared data without manually managing semaphores, reducing bugs and complexity.

Ques4:- For the Reader-writer problem, explain how starvation can occur and describe one method to prevent it.

Ans:- Starvation :- If the system gives priority to readers, a continuous stream of readers can keep arriving, causing a waiting writer to never get access → writer starvation.

Prevention Method :- use writer preference : Once a writer requests the lock, no new readers are allowed to enter. Existing readers finish, then the writer is served. This ensures no indefinite waiting.

Ques5:- In deadlock prevention, the "Hold and wait" condition can be eliminated. Explain one practical drawback of doing so in an OS.

Ans:- To prevent deadlock, processes must request all resources at once before starting.

→ Practical drawback :-

This leads to low resource utilization. A process may hold several resources it is not currently using, while others wait unnecessarily. This increases waiting time and reduces system throughput.

## Ques 6:- Distributed Deadlock Detection Simulation:

Three sites  $S_1$ ,  $S_2$  and  $S_3$  have the following wait-for graph fragments:

- $S_1 : P_1 \rightarrow P_2, P_3 \rightarrow P_4$
- $S_2 : P_2 \rightarrow P_5, P_5 \rightarrow P_6$
- $S_3 : P_6 \rightarrow P_1$

- a) Combine these fragments to form the global wait-for graph.  
b) Detect if a deadlock exists, and list the processes involved.  
c) Suggest one distributed algorithm that could be applied to detect it.

Ans. (a) Combining all edges from the three edge fragments.

Graphically:

$$P_1 \rightarrow P_2 \rightarrow P_5 \rightarrow P_6$$

and separately  $P_3 \rightarrow P_4$

(b) Detect deadlock and list processes involved

- A deadlock exists if there is a cycle in the wait-for-graph
- we have the cycle  $P_1 \rightarrow P_2 \rightarrow P_5 \rightarrow P_6 \rightarrow P_1$ . Process involved in the deadlock  $\rightarrow P_1, P_2, P_5, P_6$

(c) Suggest one distributed deadlock detection.

- Chandy - Misra - Hass (Probe) algorithm is standard choice for distributed deadlock detection

# How it works:-

- A process that suspects it is blocked (or a site periodically) initiates probe messages of from.

Ques 7 :- Distributed file system performance:

A DFS has the following characteristics:

- Average local file access time : 5ms
- Average remote file access time : 25ms
- Probability of file being remote : 0.3

- a) calculate the expected file access time.
- b) suggest one caching strategy to improve performance, and justify your choice.

Ans. Given :

- Local access time  $T_L = 5\text{ms}$
- Remote access time  $T_R = 25\text{ms}$
- Probability of file is remote  $P \rightarrow 0.3$

- a) Expected time  $E[T]$  is :

$$E[T] = P \cdot T_R + (1-P) \cdot T_L$$

Plug numbers :

$$\begin{aligned} E[T] &= 0.3 \times 25 + 0.7 \times 5 = 7.5 + 3.5 \\ &= 11.0\text{ms}. \end{aligned}$$

- b) Caching strategy suggestion .

→ Suggested strategy :- client site LRU caching with write-back for read-heavy file and validation TTL.

# justification :-

- LRU (Least Recently used) :- works well in practice because file access pattern often has temporal locality or recently used file are likely to be reused.

Ans 8:- In a concurrent system, a full checkpoint takes 200ms, while an incremental checkpoint takes 50ms. The system must maintain recovery capability within a 1-second recovery point objective (RPO).

- Propose an optimal mix of checkpointing methods over a 10-second period to meet the RPO with minimal overhead.
- Explain your reasoning.

Ans: a.) Proposed optimal mix

Steps:-

- Take one full checkpoint every 10s.
- Take incremental checkpoint every 1s between full checkpoints.

# Total Checkpoint overhead periods :-

- FULLS :  $1 \times 200 \text{ ms} = 200 \text{ ms}$
- INCREMENTALS :  $9 \times 50 \text{ ms} = 450 \text{ ms}$

$$\text{Total} = 650 \text{ ms per 10s}$$

→ Average Overhead :- 65 ms/s

b.) Reasoning :-

- RPO Constraint: The system must be prepared to restore a state no older than 1s; therefore an incremental checkpoint must be taken at least once per second.
- Full checkpoint are expensive (200ms). Doing them less frequently reduces heavy overhead full checkpoint reduce recovery time because only one full.
- Trade off: Frequent incremental add modest overhead but keep RPO tight Periodic fulls keeps recovery efficient.

Ques 9: Case Study - Global E-commerce Platform:  
A global e-commerce company runs its services using a distributed operating system deployed across multiple continents.

a.) Key challenges:

- 1.) Sudden bursty traffic :- Requests spike orders of magnitude within seconds.
- 2.) Geographic latency & data locality:- User should be serviced from the nearest region when possible.
- 3.) Hot-spots / skewed load :- Certain products / regions get disproportionate attention.
- 4.) Fairness & SLA guarantees:- Mission-critical request (Payment) must get prioritized.

b.) Fault-tolerance strategy :

1. Use multi-region deployment.
2. Replicate data based on RPO Requirements.
3. Enable automated failover to meet RTO.
4. Maintain cross-region backups & regular DR Drills.